

CS 211 S25 (Sections 5–8): Final Exam Practice

Prof. Minesh Patel

Friday, 5 May

Please sign your name and NetID below, or your work may not be graded.

I certify that the answers to this exam represent my own work and follow the RU academic integrity policies (cs.rutgers.edu/academics/undergraduate/academic-integrity-policy).

PRINT Your Name:

Your NetID (e.g., mp2099)

Instructions

- Each question is marked with its point value, totaling 400 across the exam. Choose wisely.
- Write your **final answer** inside the answer box; your work outside it.
- The Documentation appendix provides any documentation you may need (e.g., ASCII table).
- Please wait at your desk when you're finished: we will collect your exams.

Default Assumptions

Unless otherwise stated, assume:

- All provided code is syntactically correct.
- All provided code uses correct header file includes (omitted for brevity).
- All machines are 64-bit little-endian, just like your ilab environment.
- `sizeof(char) == 1`, `sizeof(short) == 2`, `sizeof(int) == 4`, `sizeof(long) == 8`
- Floating point numbers use the standard IEEE 754 32-bit representation we learned in class.
- All multiple select questions have partial credit with a hidden partial penalty for incorrect selections.

Exam Policies

- This is a closed book, closed notes exam.
- No electronic devices are permitted. Please turn them off: accessing any prohibited materials during the exam will lead to a score of 0.

Q1 [0 Points] True or False

Circle True (T) or False (F) for each of the following questions.

- T ☐ F I forgot to write down my name and ID number.
- T ☐ F Three AND gates can be used to express a XOR operation.
- ☐ T F 32-bit IEEE 754 floating-point numbers can precisely represent all 16-bit unsigned integers.
- T ☐ F The target of a RISC-V branch instruction can be located at most ± 2 MiB away from the branch instruction.
- T ☐ F UTF-8 characters are each 1-byte wide.
- ☐ T F The ISA defines the interface between the software and the hardware.
- T ☐ F Larger caches generally require less time to access.
- ☐ T F C's goto statements are comparable to unconditional jump operations in assembly.
- ☐ T F Labels in C identify specific positions in the code.

Q2 [0 Points] Boolean Algebra Simplification

Fully simplify the given boolean algebra expressions:

(a) $AB + A'B + AB'$

$B + AB' \text{ or } A + B$

(b) $(A + B + C)(A' + B)(B' + C')$

$BC' + A'B'C$

Q3 [0 Points] RISC-V Array

Suppose we run the following RISC-V assembly program:

```
1  .globl arr_func
2  arr_func:
3      mv    a2, x0
4      mv    a3, x0
5      mv    a4, x0
6      mv    a5, x0
7      mv    a6, x0
8  loop:
9      beq    a3, a1, done
10     add    a4, a3, a3
11     add    a4, a4, a4
12     add    a5, a0, a4
13     lw     a6, 0(a5)
14     blt    a6, x0, skip
15     add    a2, a2, a6
16  skip:
17     addi   a3, a3, 1
18     j      loop
19  done:
20     mv     a0, a2
21     ret
```

If this function was called so that the first input was the memory address of an `int32_t` array `[-2, 3, 7, -1, 2]`, and the second input was 5 (the length of the array), what does `arr_func` return?

Answer:

12

Q4 [0 Points] Cache Reasoning

What is the general motivation behind using caches?

- Creating the illusion that memory is both large and fast (by storing the most frequently used information where it can be accessed at higher speeds).
- Accessing main memory is significantly slower than accessing caches.

Q5 [0 Points] Struct Types

Kevin is writing a function to insert a new node at the head of a singly-linked list. Help him figure out how to update the head pointer.

```
1  #include <stdlib.h>
2
3  struct node
4  {
5      void *data;
6      struct node *next;
7  };
8
9  void insert_at_head(struct node **head)
10 {
11     struct node *new_head = (struct node *)malloc(sizeof(struct node));
12     new_head->data = NULL;
13     new_head->next = *head;
14     *head = new_head;
15 }
16
17 int main(int argc, char **argv)
18 {
19     struct node head = {NULL, NULL};
20     struct node *list = &head;
21     insert_at_head(&list);
22     return EXIT_SUCCESS;
23 }
```

Q6 [0 Points] MUX Implementation

Implement the following equation using only 2:1 MUXes. Each mux input may only be (i) a single variable (or its negation, e.g., a'); (ii) hardwired 0 or 1; (iii) or the output of another MUX. *Hint: use the MUX truth table to figure out how to perform AND and OR operations on its inputs.*

$$(a + b)cd$$

3 MUXes. We can rewrite this equation as $(ab' + 1b)(0d' + cd) = XY$, where $X = (ab' + 1b)$ and $Y = (0d' + cd)$. Then, the final mux is $0Y' + XY$.

Q7 [0 Points] Sign Extension

All of the following are well-defined typecasts in C because the destination type can correctly represent the source type.

Give the **hexadecimal representation** of the **entire destination object**, including leading zeroes. For example, a `uint16_t` storing the value 0 has representation `0x0000`.

- `int32_t a = (int8_t)-128; // a=0x`

- `uint16_t a = (int8_t)13; // a=0x`

- `int16_t a = (uint8_t)127; // a=0x`

Documentation

Selected Boolean Identities

Name	Identity	
Commutative (AND)	$a \cdot b$	$b \cdot a$
Commutative (OR)	$a + b$	$b + a$
Associative (AND)	$a \cdot (b \cdot c)$	$(a \cdot b) \cdot c$
Associative (OR)	$a + (b + c)$	$(a + b) + c$
Distributive (AND)	$a \cdot (b + c)$	$(a \cdot b) + (a \cdot c)$
Distributive (OR)	$a + (b \cdot c)$	$(a + b) \cdot (a + c)$
DeMorgan's (AND)	$(a \cdot b)'$	$a' + b'$
DeMorgan's (OR)	$(a + b)'$	$a' \cdot b'$
Absorption	$a + (a \cdot b)$	a
Absorption	$a \cdot (a + b)$	a
Redundancy	$a \cdot b' + b$	$a + b$
Redundancy	$(a + b) \cdot (a + b')$	a
Redundancy	$(a \cdot b) + (a \cdot b')$	a
Redundancy	$(a + b) \cdot (a' + b)$	b
Consensus	$a \cdot b + a' \cdot c + b \cdot c$	$a \cdot b + a' \cdot c$
Covering	$a + a' \cdot b$	$a + b$

RISC-V Reference Sheet

1. RV32I and RV64I Instructions

Inst	Name	FMT	Opcode	funct3	funct7	Description	Note
add	ADD	R	0110011	0x0	0x00	rd = rs1 + rs2	
sub	SUB	R	0110011	0x0	0x20	rd = rs1 - rs2	
xor	XOR	R	0110011	0x4	0x00	rd = rs1 ^ rs2	
or	OR	R	0110011	0x6	0x00	rd = rs1 rs2	
and	AND	R	0110011	0x7	0x00	rd = rs1 & rs2	
sll	Shift Left Logical	R	0110011	0x1	0x00	rd = rs1 << rs2	
srl	Shift Right Logical	R	0110011	0x5	0x00	rd = rs1 >> rs2	
sra	Shift Right Arith	R	0110011	0x5	0x20	rd = rs1 >> rs2	sext
slt	Set Less Than	R	0110011	0x2	0x00	rd = (rs1 < rs2)?1:0	signed
sltu	Set Less Than (U)	R	0110011	0x3	0x00	rd = (rs1 < rs2)?1:0	unsigned
addi	ADD (Immediate)	I	0010011	0x0		rd = rs1 + imm	
xori	XOR (Immediate)	I	0010011	0x4		rd = rs1 ^ imm	
ori	OR (Immediate)	I	0010011	0x6		rd = rs1 imm	
andi	AND (Immediate)	I	0010011	0x7		rd = rs1 & imm	
slli	Shift Left Logical Imm	I	0010011	0x1	imm[11:5]=0x00	rd = rs1 << imm[0:4]	
srli	Shift Right Logical Imm	I	0010011	0x5	imm[11:5]=0x00	rd = rs1 >> imm[0:4]	zext
srai	Shift Right Arith Imm	I	0010011	0x5	imm[11:5]=0x20	rd = rs1 >> imm[0:4]	sext
slti	Set Less Than Imm	I	0010011	0x2		rd = (rs1 < imm)?1:0	signed
sltiu	Set Less Than Imm (U)	I	0010011	0x3		rd = (rs1 < imm)?1:0	unsigned
lb	Load Byte	I	0000011	0x0		rd = M[rs1+imm][0:7]	sext
lh	Load Half	I	0000011	0x1		rd = M[rs1+imm][0:15]	sext
lw	Load Word	I	0000011	0x2		rd = M[rs1+imm][0:31]	sext
ld	Load Double Word	I	0000011	0x3		rd = M[rs1+imm][0:63]	
lbu	Load Byte (U)	I	0000011	0x4		rd = M[rs1+imm][0:7]	zext
lhu	Load Half (U)	I	0000011	0x5		rd = M[rs1+imm][0:15]	zext
sb	Store Byte	S	0100011	0x0		M[rs1+imm][0:7] = rs2[0:7]	
sh	Store Half	S	0100011	0x1		M[rs1+imm][0:15] = rs2[0:15]	
sw	Store Word	S	0100011	0x2		M[rs1+imm][0:31] = rs2[0:31]	
sd	Store Double Word	S	0100011	0x3		M[rs1+imm][0:63] = rs2[0:63]	
beq	Branch ==	B	1100011	0x0		if(rs1 == rs2) PC += imm	
bne	Branch !=	B	1100011	0x1		if(rs1 != rs2) PC += imm	
blt	Branch <	B	1100011	0x4		if(rs1 < rs2) PC += imm	signed
bge	Branch ≥	B	1100011	0x5		if(rs1 ≥ rs2) PC += imm	signed
bltu	Branch < (U)	B	1100011	0x6		if(rs1 < rs2) PC += imm	unsigned
bgeu	Branch ≥ (U)	B	1100011	0x7		if(rs1 ≥ rs2) PC += imm	unsigned
jal	Jump And Link	J	1101111			rd = PC+4; PC += imm	
jalr	Jump And Link Reg	I	1101111	0x0		rd = PC+4; PC = rs1 + imm	
lui	Load Upper Imm	U	0110111			rd = imm << 12	
auipc	Add Upper Imm to PC	U	0010111			rd = PC + (imm << 12)	

Note: w variants (e.g., addiw) operate on 32-bit quantities (not listed here).

2. RISC-V Pseudoinstructions

Pseudoinstruction	Base Instruction(s)	Meaning
nop	addi x0, x0, 0	No operation
li rd, immediate	<i>Myriad sequences</i>	Load immediate
mv rd, rs	addi rd, rs, 0	Copy register
not rd, rs	xori rd, rs, -1	One's complement
neg rd, rs	sub rd, x0, rs	Two's complement
negw rd, rs	subw rd, x0, rs	Two's complement word
sext.w rd, rs	addiw rd, rs, 0	Sign extend word
seqz rd, rs	sltiu rd, rs, 1	Set if = zero
snez rd, rs	sltu rd, x0, rs	Set if \neq zero
sltz rd, rs	slt rd, rs, x0	Set if < zero
sgtz rd, rs	slt rd, x0, rs	Set if > zero
beqz rs, offset	beq rs, x0, offset	Branch if = zero
bnez rs, offset	bne rs, x0, offset	Branch if \neq zero
blez rs, offset	bge x0, rs, offset	Branch if \leq zero
bgez rs, offset	bge rs, x0, offset	Branch if \geq zero
bltz rs, offset	blt rs, x0, offset	Branch if < zero
bgtz rs, offset	blt x0, rs, offset	Branch if > zero
bgt rs, rt, offset	blt rt, rs, offset	Branch if >
ble rs, rt, offset	bge rt, rs, offset	Branch if \leq
bgtu rs, rt, offset	bltu rt, rs, offset	Branch if >, unsigned
bleu rs, rt, offset	bgeu rt, rs, offset	Branch if \leq , unsigned
j offset	jal x0, offset	Jump
jal offset	jal x1, offset	Jump and link
jr rs	jalr x0, rs, 0	Jump register
jalr rs	jalr x1, rs, 0	Jump and link register
ret	jalr x0, x1, 0	Return from subroutine
call offset	auipc x1, offset[31:12] jalr x1, x1, offset[11:0]	Call far-away subroutine
tail offset	auipc x6, offset[31:12] jalr x0, x6, offset[11:0]	Tail call far-away subroutine

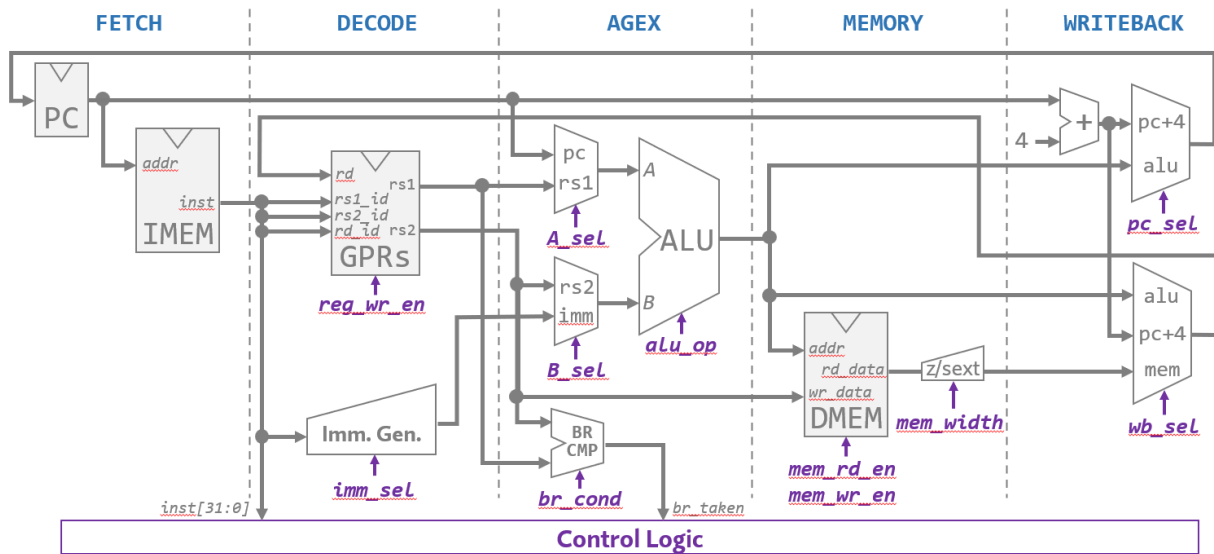
3. General-Purpose Registers

Register	ABI Name	Description	Saver
x0	zero	Zero constant	—
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	—
x4	tp	Thread pointer	—
x5-x7	t0-t2	Temporaries	Caller
x8	s0 / fp	Saved / frame pointer	Callee
x9	s1	Saved register	Callee
x10-x11	a0-a1	Fn args/return values	Caller
x12-x17	a2-a7	Fn args	Caller
x18-x27	s2-s11	Saved registers	Callee
x28-x31	t3-t6	Temporaries	Caller

4. RISC-V Instruction Encodings

31	25	24	20	19	15	14	12	11	7	6	0	
funct7		rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]				rs1		funct3		rd		opcode		I-type
imm[11:5]		rs2		rs1		funct3		imm[4:0]		opcode		S-type
imm[12-10:5]		rs2		rs1		funct3		imm[4:1-11]		opcode		B-type
		imm[31:12]						rd		opcode		U-type
		imm[20-10:1-11-19:12]						rd		opcode		J-type

Single-Cycle RV64I CPU Datapath



memset from <string.h>

```
1 void *memset( void *dest, int ch, size_t count );
```

Description:

Copies the value (unsigned char)ch into each of the first count characters of the object pointed to by dest. The behavior is undefined if access occurs beyond the end of the dest array. The behavior is undefined if dest is a NULL pointer.

Parameters:

- dest: pointer to the object to fill
- ch: fill byte
- count: number of bytes to fill
- destsz: size of the dest array

Return Value:

A copy of dest

memcpy from <string.h>

```
1 void *memcpy(void * dest, const void * source, size_t count);
```

Description:

Copies the values of count bytes from the location pointed to by source directly to the memory block pointed to by dest. To avoid overflows, the size of the arrays pointed to by both the dest and source parameters, shall be at least count bytes, and should not overlap (for overlapping memory blocks, memmove is a safer approach).

Parameters:

- dest: pointer to the object to copy to
- src: pointer to the object to copy from
- count: number of bytes to copy

Return Value:

A copy of dest

Relevant ASCII Codepoints

Hex Value	0x61	0x62	0x63	0x64	0x65	0x66	0x67	0x68	0x69	0x6a	0x6b	0x6c	0x6d
Character	a	b	c	d	e	f	g	h	i	j	k	l	m

Hex Value	0x6e	0x6f	0x70	0x71	0x72	0x73	0x74	0x75	0x76	0x77	0x78	0x79	0x7a
Character	n	o	p	q	r	s	t	u	v	w	x	y	z